

# Importing Important Libraries

```
In [ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
plt.style.use('seaborn')
```

## Reading Data Files

```
In [ ]: df = pd.read_csv('imports-85.data')
df.columns = ['symboling', 'normalized-losses', 'make', 'fuel-type', 'aspiration', 'num-o
```

```
In [ ]: df.head()
```

```
Out[ ]:
```

|   | symboling | normalized-losses | make        | fuel-type | aspiration | num-of-doors | body-style  | drive-wheels | engine-location | wheel-base |
|---|-----------|-------------------|-------------|-----------|------------|--------------|-------------|--------------|-----------------|------------|
| 0 | 3         | ?                 | alfa-romero | gas       | std        | two          | convertible | rwd          | front           | 88.6       |
| 1 | 1         | ?                 | alfa-romero | gas       | std        | two          | hatchback   | rwd          | front           | 94.5       |
| 2 | 2         | 164               | audi        | gas       | std        | four         | sedan       | fwd          | front           | 99.8       |
| 3 | 2         | 164               | audi        | gas       | std        | four         | sedan       | 4wd          | front           | 99.4       |
| 4 | 2         | ?                 | audi        | gas       | std        | two          | sedan       | fwd          | front           | 99.8       |

5 rows × 26 columns



```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 204 entries, 0 to 203
Data columns (total 26 columns):
#   Column              Non-Null Count  Dtype
---  -
0   symboling            204 non-null    int64
1   normalized-losses    204 non-null    object
2   make                 204 non-null    object
3   fuel-type            204 non-null    object
4   aspiration            204 non-null    object
5   num-of-doors         204 non-null    object
6   body-style           204 non-null    object
7   drive-wheels         204 non-null    object
8   engine-location      204 non-null    object
9   wheel-base          204 non-null    float64
10  length               204 non-null    float64
11  width                204 non-null    float64
12  height               204 non-null    float64
13  curb-weight          204 non-null    int64
14  engine-type          204 non-null    object
```

```

15 num-of-cylinders    204 non-null    object
16 engine-size        204 non-null    int64
17 fuel-system        204 non-null    object
18 bore               204 non-null    object
19 stroke             204 non-null    object
20 compression-ratio  204 non-null    float64
21 horsepower         204 non-null    object
22 peak-rpm           204 non-null    object
23 city-mpg           204 non-null    int64
24 highway-mpg        204 non-null    int64
25 price              204 non-null    object
dtypes: float64(5), int64(5), object(16)
memory usage: 41.6+ KB

```

```
In [ ]: columns_to_drop = ['bore', 'compression-ratio', 'peak-rpm', 'wheel-base', 'curb-weight',
df = df.drop(columns_to_drop, axis=1)
```

```
In [ ]: df.head()
```

```
Out[ ]:
```

|   | symboling | make        | fuel-type | aspiration | num-of-doors | body-style  | drive-wheels | engine-location | length | width | height |
|---|-----------|-------------|-----------|------------|--------------|-------------|--------------|-----------------|--------|-------|--------|
| 0 | 3         | alfa-romero | gas       | std        | two          | convertible | rwd          | front           | 168.8  | 64.1  | 48.8   |
| 1 | 1         | alfa-romero | gas       | std        | two          | hatchback   | rwd          | front           | 171.2  | 65.5  | 52.4   |
| 2 | 2         | audi        | gas       | std        | four         | sedan       | fwd          | front           | 176.6  | 66.2  | 54.3   |
| 3 | 2         | audi        | gas       | std        | four         | sedan       | 4wd          | front           | 176.6  | 66.4  | 54.3   |
| 4 | 2         | audi        | gas       | std        | two          | sedan       | fwd          | front           | 177.3  | 66.3  | 53.1   |

## Cleaning the DataFrame

```
In [ ]: df_clean = df.replace(to_replace='?', value = np.nan)
```

```
In [ ]: df_clean = df_clean.fillna(value = df_clean.median())
```

```
In [ ]: df_clean = df_clean.dropna()
```

```
In [ ]: df_clean.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 202 entries, 0 to 203
Data columns (total 20 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   symboling        202 non-null    int64
1   make             202 non-null    object
2   fuel-type        202 non-null    object
3   aspiration        202 non-null    object
4   num-of-doors     202 non-null    object

```


```
5  body-style      202 non-null  object
6  drive-wheels   202 non-null  object
7  engine-location 202 non-null  object
8  length         202 non-null  float64
9  width          202 non-null  float64
10 height         202 non-null  float64
11 engine-type    202 non-null  object
12 num-of-cylinders 202 non-null  object
13 engine-size    202 non-null  int64
14 fuel-system    202 non-null  object
15 stroke         202 non-null  object
16 horsepower     202 non-null  object
17 city-mpg       202 non-null  int64
18 highway-mpg    202 non-null  int64
19 price         202 non-null  object
dtypes: float64(3), int64(4), object(13)
memory usage: 33.1+ KB
```

```
In [ ]: df_clean['price'] = df_clean['price'].astype('float64')
```


## Visualizing Data to make a Co-relation Matrix

```
In [ ]: import seaborn as sns
```

```
In [ ]: sns.pairplot(df_clean)
plt.show()
```



```
In [ ]: sns.heatmap(df_clean.corr(),annot=True)
plt.show()
```



## Encoding Class Data to create Continuous Regression Data


```
In [ ]: from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
columns_to_encode = ['make', 'fuel-type', 'aspiration', 'num-of-doors', 'body-style', 'drive-wheels']
for i in columns_to_encode:
    df_clean[i] = le.fit_transform(df_clean[i].astype(str))
```

```
In [ ]: df_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 202 entries, 0 to 203
Data columns (total 20 columns):
#   Column                Non-Null Count  Dtype
---  -
0   symboling              202 non-null    int64
1   make                   202 non-null    int32
2   fuel-type              202 non-null    int32
3   aspiration              202 non-null    int32
4   num-of-doors           202 non-null    int32
5   body-style             202 non-null    int32
6   drive-wheels           202 non-null    int32
7   engine-location        202 non-null    int32
8   length                 202 non-null    float64
9   width                  202 non-null    float64
10  height                 202 non-null    float64
11  engine-type            202 non-null    int32
12  num-of-cylinders       202 non-null    int32
13  engine-size            202 non-null    int64
14  fuel-system            202 non-null    int32
15  stroke                 202 non-null    int32
16  horsepower             202 non-null    int32
17  city-mpg               202 non-null    int64
18  highway-mpg            202 non-null    int64
```

```
19 price                202 non-null    float64
dtypes: float64(4), int32(12), int64(4)
memory usage: 23.7 KB
```

```
In [ ]: plt.figure(figsize=(20,20))
sns.heatmap(df_clean.corr(),annot=True)
plt.show()
```



## Purging Features with very low correlation

```
In [ ]: df_clean = df_clean.drop(['symboling', 'body-style', 'num-of-doors', 'stroke', 'num-o
```

```
In [ ]: df_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 202 entries, 0 to 203
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---

```

```

0  make                202 non-null    int32
1  fuel-type           202 non-null    int32
2  aspiration           202 non-null    int32
3  drive-wheels        202 non-null    int32
4  engine-location     202 non-null    int32
5  length              202 non-null    float64
6  width               202 non-null    float64
7  height              202 non-null    float64
8  engine-size         202 non-null    int64
9  fuel-system         202 non-null    int32
10 horsepower          202 non-null    int32
11 city-mpg            202 non-null    int64
12 highway-mpg         202 non-null    int64
13 price               202 non-null    float64
dtypes: float64(4), int32(7), int64(3)
memory usage: 18.1 KB

```

```
In [ ]: X,Y = df_clean.values[:, :-1], df_clean.values[:, -1]
```

```
In [ ]: X.shape, Y.shape
```

```
Out[ ]: ((202, 13), (202,))
```

## Data Normalization

```
In [ ]: u = np.mean(X, axis=0)
# print(u.shape)
std = np.std(X, axis=0)
X = (X-u)/std
```

```
In [ ]: ones = np.ones((X.shape[0],1))
X = np.hstack((ones,X))
print(X)
```

```

[[ 1.          -1.98265203  0.32221908 ... -1.46925631 -0.64279561
  -0.54119201]
 [ 1.          -1.98265203  0.32221908 ... -0.63760179 -0.9493947
  -0.68675896]
 [ 1.          -1.82168108  0.32221908 ... -1.63558721 -0.18289698
  -0.10449113]
 ...
 [ 1.           1.3977378   0.32221908 ... -1.02570723 -1.10269424
  -1.12345984]
 [ 1.           1.3977378  -3.10347852 ... -1.58014358  0.12370211
  -0.54119201]
 [ 1.           1.3977378   0.32221908 ... -1.35836904 -0.9493947
  -0.83232592]]

```

## Splitting the Data into Test Train Split

```
In [ ]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.1, random_stat
```

## Multiple Linear Regression

```
In [ ]: def r2_score(y,y_):
    num = np.sum((y-y_)**2)
    denom = np.sum((y-y.mean())**2)
    score = 1-(num/denom)
    return score*100

def train_val_data(X,Y):
    train_X,val_X,train_Y,val_Y = train_test_split(X, Y, test_size=0.05,shuffle=True)
    return train_X,val_X,train_Y,val_Y
```

```
In [ ]: def hypothesis(X,theta):
    return np.dot(X,theta)

def error(X,y,theta):
    e = 0.0
    m = X.shape[0]
    y_ = hypothesis(X,theta)
    e = (np.sum((y-y_)**2))
    return e/m

def gradient(X,y,theta):
    y_ = hypothesis(X,theta)
    grad = np.dot(X.T,(y_-y))
    m = X.shape[0]
    return grad/m

def gradient_Descent(X,y,learning_rate = 0.1,max_iters = 300):
    n = X.shape[1]
    theta = np.zeros((n,))
    error_list = []
    val_score = []

    for i in range(max_iters):
        train_X,val_X,train_Y,val_Y = train_val_data(X,y)

        e = error(train_X,train_Y,theta)
        error_list.append(e)


        grad = gradient(train_X,train_Y,theta)
        theta = theta - learning_rate*grad

        y_ = hypothesis(val_X,theta)
        val_score.append(r2_score(val_Y,y_))
    return theta,error_list,val_score
```

```
In [ ]: theta,error_list,val_score = gradient_Descent(X_train,y_train)
```

```
In [ ]: plt.plot(error_list)
plt.show()
```





```
In [ ]: val_score = pd.DataFrame(np.array(val_score),columns=["Val Score"])
print(val_score.describe())
```

|       | Val Score    |
|-------|--------------|
| count | 300.000000   |
| mean  | 53.485459    |
| std   | 111.013255   |
| min   | -1539.427710 |
| 25%   | 55.857035    |
| 50%   | 78.884615    |
| 75%   | 89.711316    |
| max   | 98.182104    |

```
In [ ]: y_ = hypothesis(X_test,theta)
print("Overall R2-Score: ",r2_score(y_test,y_))
```

Overall R2-Score: 93.73951859962185

## Ridge Regression

```
In [ ]: def hypothesis_r(X,theta):
        return np.dot(X,theta)

def error_r(X,y,theta, lam):
    e = 0.0
    m = X.shape[0]
    y_ = hypothesis_r(X,theta)
    e = np.sum((y-y_)**2) + lam*np.sum(np.square(theta))
    return e/m

def gradient_r(X,y,theta,lam):
    y_ = hypothesis_r(X,theta)
    grad = np.dot(X.T,(y-y_)) + lam*theta
    m = X.shape[0]
    return grad/m

def gradient_Descent_r(X,y,lam,learning_rate = 0.1,max_iters = 300):
    n = X.shape[1]
    theta = np.zeros((n,))
    error_list = []
    val_score = []
```

```

for i in range(max_iters):
    train_X, val_X, train_Y, val_Y = train_val_data(X, y)

    e = error_r(train_X, train_Y, theta, lam)
    error_list.append(e)

    grad = gradient_r(train_X, train_Y, theta, lam)
    theta = theta - learning_rate * grad

    y_ = hypothesis_r(val_X, theta)
    val_score.append(r2_score(val_Y, y_))
return theta, error_list, val_score

```

```


In [ ]: theta, error_list, val_score = gradient_Descent_r(X_train, y_train, lam=10)

```

```

In [ ]: plt.plot(error_list)
plt.title("Ridge Regression Error")
plt.show()

```



```

In [ ]: val_score = pd.DataFrame(np.array(val_score), columns=["Val Score Ridge Regression"])
print(val_score.describe())

```

```

      Val Score Ridge Regression
count          300.000000
mean             63.269598
std             65.456889
min            -643.204399
25%             65.353973
50%             82.525613
75%             89.126761
max             97.744985

```

```

In [ ]: y_ = hypothesis_r(X_test, theta)
print("Overall R2-Score: ", r2_score(y_test, y_))

```

```

Overall R2-Score: 94.26890687581202

```

# Lasso Regression

In [ ]:

```
def hypothesis_l(X,theta):
    return np.dot(X,theta)

def error_l(X,y,theta, lam):
    e = 0.0
    m = X.shape[0]
    y_ = hypothesis_l(X,theta)
    e = np.sum((y-y_)**2) + lam*np.sum(np.abs(theta))
    return e/m

def gradient_l(X,y,theta,lam):
    y_ = hypothesis_l(X,theta)
    grad = np.dot(X.T,(y-y_)) + lam*(2*np.sinc(theta)-1)
    m = X.shape[0]
    return grad/m

def gradient_Descent_l(X,y,lam,learning_rate = 0.1,max_iters = 300):
    n = X.shape[1]
    theta = np.zeros((n,))
    error_list = []
    val_score = []

    for i in range(max_iters):
        train_X,val_X,train_Y,val_Y = train_val_data(X,y)

        e = error_l(train_X,train_Y,theta,lam)
        error_list.append(e)

        grad = gradient_l(train_X,train_Y,theta,lam)
        theta = theta - learning_rate*grad


        y_ = hypothesis_l(val_X,theta)
        val_score.append(r2_score(val_Y,y_))
    return theta,error_list,val_score
```

In [ ]:

```
theta,error_list,val_score = gradient_Descent_l(X_train,y_train,lam=10)
```

In [ ]:

```
plt.plot(error_list)
plt.title("Lasso Regression Error")
plt.show()
```



```
In [ ]: val_score = pd.DataFrame(np.array(val_score), columns=["Val Score Lasso Regression"])
        print(val_score.describe())
```

|       | Val Score Lasso Regression |
|-------|----------------------------|
| count | 300.000000                 |
| mean  | 65.963418                  |
| std   | 47.818503                  |
| min   | -281.041242                |
| 25%   | 63.649129                  |
| 50%   | 81.367128                  |
| 75%   | 90.137017                  |
| max   | 99.064463                  |

```
In [ ]: y_ = hypothesis_r(X_test, theta)
        print("Overall R2-Score: ", r2_score(y_test, y_))
```

Overall R2-Score: 93.77253605140649

```
In [ ]:
```